TECH & ACCOUNTANCY CAREER FAIR 2026

30 Days to Agentic AI Employability

Learn the foundations. Build a practical workflow. Show employers what you can do.

Beginner-friendly · 30–45 minutes a day

Nexius Academy
Practical Agentic AI learning

START HERE

What this 30-day plan will help you do

You will not become an Agentic AI expert in 30 days. You can, however, build a credible foundation and one piece of work you can explain to an employer.

Understand

Explain how Agentic AI differs from ordinary chat, what agents need to work and where their limits are.

Design

Turn one bounded workplace task into steps, tools, checks, exceptions and clear human responsibilities.

Test

Run normal and failure cases using public or synthetic information. Record what worked, failed and changed.

Show

Produce a portfolio workflow, evidence-backed CV bullet and 60-second interview explanation.

How to use the workbook

- ☐ Set aside 30–45 minutes on most days. Allow more time for the Week 3 project.
- ☐ Complete the small output each day. The outputs combine into your final evidence pack.
- ☐ Use public or synthetic information. Do not upload confidential employer, customer or personal data.
- ☐ Write in your own words. Employers need to see your judgment, not polished AI language.

By Day 30
You should be able to show one tested workflow, one portfolio case study, one CV bullet and one interview explanation.

Want a guided start?
Join the 16-hour Agentic AI Foundations course and practise these concepts with trainers, tools and workplace examples.

COURSE DETAILS · PAGE 16

Four stages. One employability outcome.

The stages happen in this order because useful AI work starts with understanding, moves into design, survives testing and ends with evidence.

01
Days 1–7
Understand

Learn the language, building blocks, limits and workplace relevance of Agentic AI.

Checkpoint: Explain Agentic AI in your own words.

02
Days 8–14
Design

Choose one task and map its goal, inputs, steps, tools, checks and owner.

Checkpoint: Workflow blueprint v1.

03
Days 15–21
Build and test

Use safe data, test normal and failure cases, then improve the workflow.

Checkpoint: Tested workflow v2.

04
Days 22–30
Show

Package your work for a portfolio, CV and interview conversation.

Checkpoint: Employability pack.

> **“I used AI” is weak evidence. “I designed the workflow, limited the data, added review gates, tested failure cases and documented the result” is credible.**

The five capabilities you will practise

Problem framing
Define the work before selecting a tool.

Workflow design
Turn work into visible steps and decisions.

AI collaboration
Provide context and improve through feedback.

Verification
Check evidence, privacy, policy and exceptions.

Application evidence
Show what you built, tested and owned.

Prefer to learn with a practitioner?
The 16-hour Agentic AI Foundations course covers agents, context, tools, workflows, safety and human review.

COURSE DETAILS · PAGE 16

Days 1–3: Learn what makes AI agentic

1

AI, Generative AI and Agentic AI

Learn: Traditional automation follows fixed rules. Generative AI creates content. Agentic AI works towards a goal through steps, tools, checks and hand-offs.

Practise: Compare a spreadsheet formula, a chatbot drafting an email and an AI workflow that checks sources before producing a report.

Produce: A three-column comparison.

2

AI changes tasks before titles

Learn: A job contains many tasks. Some may be prepared or accelerated by AI while judgment and accountability remain human.

Practise: Save five job advertisements for a role you want. Highlight repeated tasks.

Produce: Five recurring tasks from your target role.

3

Recognise the building blocks

Learn: A useful agentic workflow has a goal, context, instructions, trusted inputs, tools, steps, checks, exceptions and an accountable human.

Practise: Take a familiar task such as preparing a report. Label its goal, information sources, working steps, checks and final owner.

Produce: One annotated workflow example.

Starter question

What must be true before this work can begin?

Starter question

What could go wrong if the AI guesses?

Starter question

Who remains responsible for the result?

If you are stuck

Use a familiar task: summarising meeting notes, comparing options or preparing a weekly update. You are learning the anatomy of the work, not building anything yet.

COURSE DETAILS · PAGE 16

Days 4–7: Context, tools and human control

4

Goals, context and instructions

Learn: The agent needs a useful outcome, relevant background, trusted information, constraints and an output format.

Practise: Rewrite “Summarise this” to include the user, purpose, source boundary, checks and format.

Produce: Your first context brief.

5

Tools and actions

Learn: Agents may read files, search approved sources, calculate, create drafts or update records. Tool access is not decision authority.

Practise: Match five tasks to appropriate tools and identify one action that still needs approval.

Produce: Task-and-tool map.

6

Limits and safety

Learn: AI can be wrong, overconfident, incomplete or unsafe with sensitive information.

Practise: Sort actions into “AI may prepare”, “human must review” and “human only”.

Produce: Human-control checklist.

7

Explain it in one minute

Learn: Clear explanation is evidence of understanding.

Practise: Explain Agentic AI to a hiring manager without jargon.

Produce: “What Agentic AI means for my role” summary.

Week 1 quality check
You can explain the difference between chat and an agentic workflow, name its building blocks, identify a suitable task and state which decisions remain human.

Build the foundation with practitioners.
The 16-hour course covers agent concepts, tools, context, safe review and practical workplace use.

COURSE DETAILS · PAGE 16

Checkpoint: explain Agentic AI in your own words

Use this page to turn what you learned into a short explanation you can use in a conversation or interview.

1. In ordinary workplace language, Agentic AI is:

2. It differs from an ordinary chatbot because:

3. In my target role, it could assist with:

4. The person must remain responsible for:

5. One risk I would design for is:

What good looks like

Your explanation mentions a goal, multiple steps, tools or information, checks, exceptions and a responsible person.

Avoid

Do not describe Agentic AI as “AI that does everything automatically”. That removes the controls employers need to trust the work.

COURSE DETAILS · PAGE 16

Days 8–10: Choose and frame the project

8

Choose one bounded task

Learn: A good first project is useful, repeatable, safe to test and easy for you to verify.

Practise: Score three tasks from 1–5 for role relevance, data safety, repeatability and ease of checking.

Produce: Selected project task.

9

Define the outcome

Learn: A workflow needs a user, purpose, expected output and success criteria.

Practise: Complete: “This workflow helps ___ complete ___ by producing ___ while keeping ___ under human control.”

Produce: One-sentence project statement.

10

Identify trusted inputs

Learn: AI output cannot be more trustworthy than the information allowed into the workflow.

Practise: List every required source. Label it public, synthetic, personal, confidential or restricted. Remove anything unsuitable for practice.

Produce: Approved-input list and data boundary.

Task-selection test

Question	Good sign	Warning sign
Can you explain the task?	You understand the normal process.	You cannot judge whether the output is correct.
Can you test it safely?	Public or synthetic data is sufficient.	It requires real customer or employer data.
Can you keep it bounded?	One clear output and owner.	It attempts to automate an entire role.

If you are stuck
Use one of the Week 3 projects: Tech support-ticket triage or Accountancy variance analysis. Both include clear inputs, outputs and human review.

Choose a project with trainer guidance.
The 16-hour course helps you identify suitable use cases and avoid unsafe or overambitious first projects.

COURSE DETAILS · PAGE 16

Days 11–14: Build the workflow blueprint

11

Break the work into steps

Map the process from input to output. Include decisions, hand-offs, missing-information requests and final approval.

Produce: Process map v1.

12

Assign responsibilities

Mark each step as AI, human, conventional software or shared. Keep judgment and authority explicit.

Produce: Responsibility table.

13

Add checks and exceptions

Define what happens when information is missing, sources conflict, data is sensitive, confidence is low or a tool fails.

Produce: Review and exception rules.

14

Complete the blueprint

Combine the goal, inputs, steps, tools, responsibilities, checks, stop conditions and owner.

Produce: Workflow blueprint v1.

Approved inputs
What may enter the workflow?

→

AI-assisted work
What can the agent prepare or execute?

→

Human decision
Who checks, approves or escalates?

Week 2 quality check
A stranger can see what the workflow does, what information it uses, what the AI handles, what gets checked and when the process must stop.

Design your first governed workflow.
In the 16-hour course, learners practise context engineering, workflow design, tools, reusable skills and human review.

COURSE DETAILS · PAGE 16

Workflow blueprint template

Goal and user

What should the workflow accomplish, and who uses the result?

Expected output

Approved inputs

List only trusted, permitted sources.

Excluded data

AI-assisted steps

Tools required

Human-owned steps

Final approver

Checks before approval

Stop and escalation conditions

Starter questions
What evidence supports the output? What would make the result unsafe to use? What decision cannot be delegated? What should the workflow record for later review?

Turn this blueprint into a working practice exercise.
The 16-hour Agentic AI Foundations course provides guided, no-code practice with agents, tools and workflow controls.

COURSE DETAILS · PAGE 16

Choose a Tech or Accountancy practice project

Both projects teach the same agentic pattern: approved inputs, structured preparation, evidence checks, exception handling and a human decision.

Tech: Support-ticket triage

Scenario: Use fictional support requests, a simple knowledge base and escalation rules.

AI may assist with: extracting symptoms, classifying requests, finding approved guidance, identifying missing information and drafting a hand-off.

Human owns: severity, security, access decisions, technical validation, escalation and customer communication.

Portfolio evidence: Triage table, routing rules, missing-information checklist, draft hand-off and correction log.

Accountancy: Variance commentary

Scenario: Use a synthetic budget-versus-actual table, explanatory notes and materiality rules.

AI may assist with: checking completeness, identifying unusual movements, grouping exceptions and drafting commentary.

Human owns: source accuracy, materiality, accounting interpretation, investigation and final sign-off.

Portfolio evidence: Validated table, exception list, commentary, source traceability and maker-checker log.

Use the same project boundaries

- ☐ Use public or synthetic information only.
- ☐ Do not connect the exercise to a real customer, support or accounting system.
- ☐ Do not let the AI approve, send, close or post anything.
- ☐ Record every material correction and human decision.

Project goal
You are not trying to prove that AI can replace the role. You are proving that you can design, test and explain a controlled AI-assisted workflow.

Build a project with guided practice.
The 16-hour course helps beginners work with agent concepts, no-code tools, reusable skills and safer review habits.

COURSE DETAILS · PAGE 16

Days 15–17: Prepare the first working version

15 Understand the project

Review the scenario, inputs, expected output and human responsibilities. Remove anything outside the project boundary.

Produce: Completed project brief.

16 Prepare safe test data

Create realistic synthetic examples. Include enough variation to reveal weaknesses.

Produce: Test-data pack.

17 Write agent instructions

State the goal, context, sources, steps, constraints, format, checks and stop conditions.

Produce: Instruction set v1.

Instruction builder

Component	Starter wording
Role and task	"Assist with [bounded task]. You prepare; the named human reviews and decides."
Approved sources	"Use only the supplied information. Do not invent missing facts."
Working steps	"Extract, classify, check, flag exceptions and prepare the required output."
Stop conditions	"Stop and request review when information is missing, conflicting, sensitive or outside scope."
Output	"Return a structured table, evidence links, unresolved questions and draft output."

If you are stuck
Write the process as if you were briefing a careful new colleague. Tell them what they may use, what to do, what not to assume and when to ask for help.

Days 18–21: Test failure, correct and retest

18

Run a normal case

Use a complete, straightforward example. Save the input, output and every manual intervention.

Produce: Normal test result.

19

Verify the result

Trace claims to sources, recalculate figures, check completeness and review the human hand-off.

Produce: Review checklist and correction log.

20

Run an edge case

Use missing information, conflicting sources, a sensitive request, a wrong formula or an out-of-scope instruction.

Produce: Failure-case result.

21

Improve and retest

Fix the workflow, not only the wording. Add missing checks, reorder steps and strengthen stop conditions.

Produce: Workflow v2 and final test log.

Tech edge cases

- ☐ Suspected phishing request
- ☐ Missing device or error details
- ☐ Conflicting urgency information
- ☐ Access change needing approval

Accountancy edge cases

- ☐ Incorrect formula
- ☐ Missing explanation
- ☐ Large percentage, low dollar impact
- ☐ Conflicting management notes

Week 3 quality check
You can show what worked, what failed, what you corrected, what remained human and whether the revised workflow handled the same problem better.

Learn to test, review and improve safely.
The 16-hour course includes guided practice in AI limitations, review habits, workflow risks and human-in-the-loop controls.

COURSE DETAILS · PAGE 16

WEEK 3 WORKSHEET

Test and correction log

Test	What you changed	Expected result	What happened	Correction
Normal case 1				
Normal case 2				
Edge case				
Retest				

What still needs human judgment?

What evidence proves the final result?

Good enough to continue
The workflow handles the normal case, stops or escalates safely in the edge case, and leaves a clear record of sources, checks and human decisions.

COURSE DETAILS · PAGE 16

Days 22–25: Build the portfolio case

22 Organise the evidence

Collect the blueprint, instructions, test inputs, normal and edge cases, checklist, corrections and final output.

Produce: Evidence folder.

23 Compare before and after

Use honest measures: time, completeness, consistency, errors caught, turnaround or visibility of checks.

Produce: Comparison table.

24 Write the case study

Cover the problem, original process, workflow design, controls, tests, result and limitations.

Produce: One-page case study.

25 Create the visual

Show inputs, AI-assisted steps, checks, exceptions, human decision and final output.

Produce: Workflow diagram.

One-page case-study structure

Section	What to write
Problem	What work needed improvement, and who needed the result?
Design	What role did AI play, and what remained human?
Controls	How did you protect data, verify output and handle exceptions?
Evidence	What did you test, what failed and what changed?
Result	What improved, what did not and what would you do next?

Avoid inflated claims

Do not say “automated the process” if a person still checks, decides or approves. Describe the workflow accurately: AI-assisted, human-reviewed and tested.

Turn learning into visible workplace capability.

The 16-hour course gives you practical material and guided exercises you can continue developing into portfolio evidence.

COURSE DETAILS · PAGE 16

Days 26–30: Use the work in your job search

26

Write the CV bullet

Use: “Designed and tested an AI-assisted workflow for [task], incorporating [controls] and improving [honest result].”

Produce: One CV bullet.

27

Write a professional summary

Explain what you built, why it mattered, how you tested it and what remained under human control.

Produce: 100-word summary.

28

Prepare the interview answer

Cover the task, workflow, AI role, checks, failure, correction and your responsibility.

Produce: 60-second answer.

29

Get one useful critique

Ask a trainer, practitioner, mentor or peer to challenge the workflow and evidence.

Produce: Feedback record.

30

Complete the employability pack

Finalise the tested workflow, one-page case study, visual, CV bullet, professional summary, interview explanation and reflection on limitations.

Produce: Your Agentic AI employability pack.

Interview line
“I used AI to support a bounded work task. I mapped the process, limited the data, tested normal and failure cases, added human review and documented what changed.”

COURSE DETAILS · PAGE 16

YOUR NEXT STEP

Learn Agentic AI with guided, hands-on practice

Agentic AI Foundations for Non-Technical Professionals is a practical, no-code, 16-hour course for turning everyday work into governed AI-assisted workflows with human review built in.

Understand agents

Learn how goals, context, tools and multi-step work differ from ordinary chat.

Build workflows

Practise turning workplace tasks into roles, steps, checks and reusable agent skills.

Keep control

Review AI output, protect sensitive information and keep human judgment explicit.

SEPTEMBER COHORT

18 & 25 September 2026

9:00am–6:00pm
In person · Singapore Institute of Management
461 Clementi Road

OCTOBER COHORT

9 & 16 October 2026

9:00am–5:00pm
In person · Singapore Institute of Management
461 Clementi Road

Scan to view course and registration details

Use the course as a guided route through the learning plan. Bring a target role or task, practise during the course, then continue testing and documenting your project afterwards.

SkillsFuture-eligible. Final payable amount and live seat availability depend on learner eligibility and the registration channel.

Do it once. Skill it up. Do it again.

Source: Nexius Academy AI Career Readiness landing page and Agentic AI course registration destination, reviewed 29 August 2026 UTC.

Ready for guided practice?

Scan the QR code and view the next available 16-hour Agentic AI Foundations class.

SCAN TO VIEW COURSE